Report: 640 vs 1280 vs tiling on VisDrone2019-DET

Report of project **P17 Small-Object Detection in Aerial Images**: "*Compare standard resizing with tiling or higher-resolution inference. Report AP/recall, runtime, and failure cases involving very small objects.*"

All numbers are on VisDrone2019-DET **test-dev** (1,610 images, 75,101 objects), computed with the weights from the full fine-tune (50 epochs) by one shared evaluator, `src/evaluation/metrics.py`. Raw numbers (0–1, every metric): `results/metrics_all.csv`; column guide: `results/README.md`. The numbers follow our protocol and are not official VisDrone leaderboard scores.

0. Summary

Question	Answer
Does higher inference resolution help small objects?	**Yes, a lot.** Both 1280 and tiling raise recall on tiny objects ($<16^2$ px) by 1.6–1.7× (25 → 41–47 %) and AP on tiny objects by 1.4–1.7×.
Which is better, 1280 or tiling?	**1280.** Higher mAP (+4 points vs +0.9), higher AP on tiny and small objects, and **2.2× faster** than tiling. Tiling only wins recall on tiny objects (+1–2 points), and AP on tiny objects only on the largest frames (~1920 px), where it actually adds pixels over 1280.
What is the cost?	Runtime per image: 1280 ×3.3–3.7 vs 640, tiling ×7.2–8.1. Large objects ($\geq 96^2$ px) **regress** in both modes (1280: −4…−10 AP, tiling: −14…−18 AP).
What matters more for small objects: resolution or model capacity?	**Resolution.** yolo26x instead of yolo26m gives +1.6–1.9 mAP at ×2.2–2.5 runtime; 1280 instead of 640 gives +3.8–4.1 mAP. yolo26m@1280 (25.4) beats yolo26x@640 (23.2).
Which tile overlap?	20 % is a sensible trade-off: 0→20 % gives +0.6 mAP, 20→40 % only +0.03 mAP at almost 2× p95 runtime.

1. What is compared

Each model is fine-tuned **once**, at `imgsz=640`. The **same weights** are then evaluated under three input modes. Only inference changes, so differences between modes are the effect of the input resolution, not of training.

Mode (task wording)	In the project	What happens to the image	Forward passes per image
standard resizing	`full-640`	the whole image is downscaled to a 640 long side (1920×1080 → 640×360)	1
higher-resolution inference	`full-1280`	the whole image is resized to a 1280 long side	1
tiling	`tiled-ov20`	the image is cut at **native resolution** into 640×640 crops with 20 % overlap; boxes are mapped back to image coordinates and merged with NMS (IoU 0.7)	2–8

Two models × three modes = **6 main experiments**, plus a tile-overlap sensitivity analysis (0 / 20 / 40 %) for yolo26m.

What "1280" means for different frames. `imgsz` is not a fixed network input (YOLO is fully convolutional and accepts any size divisible by 32); it is the length the long side is resized to. Frames smaller than `imgsz` are **upscaled** (no new detail, only interpolation). Test-dev has six native sizes:

Native size	Images	Scale at 640	Scale at 1280	Tiles (20 %)
960×540	49	×0.67	**×1.33 (upscaling)**	2
1360×765	167	×0.47	×0.94	6
1400×788	933	×0.46	×0.91	6
1400×1050	274	×0.46	×0.91	6
1916×1078	151	×0.33	×0.67	8
1920×1080	36	×0.33	×0.67	8

For **85 %** of the images (width 1360–1400), 1280 is almost native resolution (×0.91–0.94). Tiling (×1.0) adds only ~10 % per side (~20 % in area) over 1280; only for the 187 frames of width ~1920 does it add substantially more (×1.0 vs ×0.67, 2.2× in area). This explains most of the results below.

2. Data

Official VisDrone2019-DET splits, nothing reshuffled. 10 classes (pedestrian, people, bicycle, car, van, truck, tricycle, awning-tricycle, bus, motor); "ignored regions" and "don't care" areas are excluded from training targets and respected by the evaluator (predictions inside them are not penalised).

Split	Images	Objects	Role
train	6,471	343,197	training the weights
val	548	38,759	**only** best-epoch selection and early stopping
test-dev	1,610	75,101	**only** the numbers in this report

No hyperparameters were tuned on val — they were fixed by the protocol in advance; the only value selected was the batch size (by GPU memory, not by accuracy; see §3).

Test-dev objects by size (area in original-image pixels, disjoint groups):

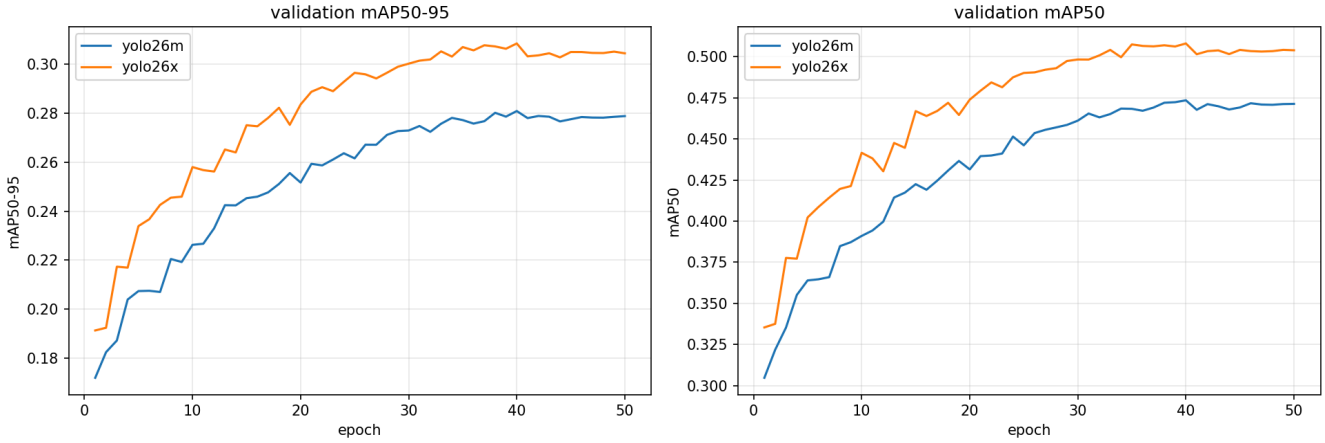
Group	Area	Objects	Share
tiny	$< 16^2$ (256 px²)	27,345	36.4 %
small	16^2–32^2	23,491	31.3 %
medium	32^2–96^2	21,856	29.1 %
large	$\geq 96^2$	2,409	3.2 %

Two thirds of all objects are small. Test-dev has more tiny objects than train (36.4 % vs 26.0 %, see the EDA), so absolute tiny-object numbers are lower than val would suggest; this does not affect the comparison of modes, which are all scored on the same test-dev.

3. Training

Both models start from COCO-pretrained weights and are fine-tuned with **one shared config** (`results/training_config.json`): `imgsz=640`, up to 50 epochs, `patience=15`, AdamW (`optimizer=auto`), `lr0=0.01` scaled to the batch (`nbs=64`), cosine LR, `warmup_epochs=3`, `weight_decay=5e-4`, default Ultralytics augmentation, `close_mosaic=10`, seed 42, all layers trainable. Batch 16 was chosen by a pilot on yolo26x (1 epoch on 2 % of train): peak memory 15.1 GB on a 16 GB T4 (batch 4: 5.4 GB, 8: 9.5 GB). Training ran on Kaggle, 2× Tesla T4, one model per GPU in parallel.

Model	Parameters	Epochs	Best epoch	val mAP50–95	val mAP50	Wall time
yolo26m	21.8 M	50	40	0.281	0.473	3.2 h
yolo26x	58.8 M	50	40	0.308	0.508	7.9 h



Both models plateau around epoch 40; early stopping did not trigger, so 50 epochs were enough. For comparison, the short 5-epoch baseline run gave test-dev mAP50–95 at 640 of 18.5 (m) and 20.7 (x); the full run gives **21.4 and 23.2** (stage 1 was scored with the older evaluator, see §4, so the comparison is approximate).

> The val metrics above come from the built-in Ultralytics validator (ignored regions not respected, different detection cap). They cannot be compared directly with the test-dev numbers below.

---

4. Evaluation protocol

All modes are scored by **one** evaluator; all parameters were fixed before looking at results:

Metric	Definition
mAP50–95	COCO AP, IoU 0.50:0.05:0.95, confidence threshold 0.001
mAP50	the same at IoU 0.5
Recall	share of objects found at IoU ≥ 0.5 **and** confidence ≥ 0.25
AP / AP50 / Recall per size	the same, restricted to objects of one group (tiny / small / medium / large)
Runtime	ms per **original** image, batch 1, 10 warm-up images, `cuda.synchronize`; tiling counts all tiles + merging; first 200 test-dev images; Tesla T4

Details that matter for a fair comparison:

1. **Same detection cap — 1000 per image — in every mode.** By default Ultralytics returns ≤300 boxes per *forward pass*: 300 per image at 640/1280, but 300 per *tile* with tiling (up to 2,400 per image). On crowded frames (up to ~900 objects) this favoured tiling. The cap is now equal (for tiling, after merging), and COCOeval also uses up to 1000 detections. The 5-epoch stage-1 numbers in `docs/baseline_results.md` and `docs/tiling_results.md` were computed *before* this fix.
2. **Size groups are disjoint** (in COCO, "small" includes tiny). COCO-small is also kept in `metrics_all.csv` (`ap_small`) for comparability.
3. **AP50 next to AP50–95.** For a 10×10 px box, a 1–2 px shift already drops IoU below 0.75, so AP50–95 on tiny objects is depressed by the metric itself. AP50 answers "was the object found?", AP50–95 "how precisely was it localised?".
4. All coordinates are in original-image pixels; ignored regions are respected.
5. Single run, no confidence intervals: treat differences below ~0.5 points as noise. The conclusions below rest on differences of several points.

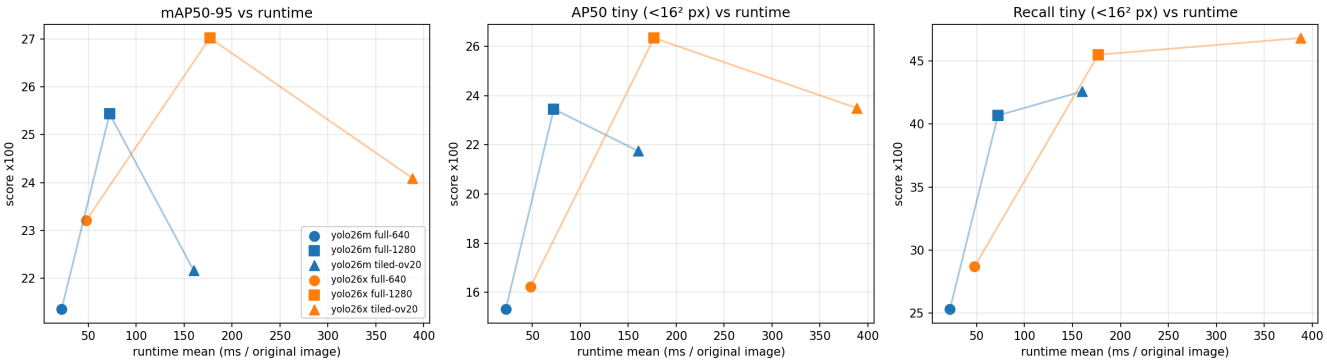
---

5. Results

All metrics below are ×100 (AP / recall points). Δ is the difference from `full-640` of the same model, in points; × is a ratio.

5.1. Overall metrics and runtime (Table 1)

Model	Mode	mAP50–95	mAP50	Recall	ms/img (mean)	ms/img (p95)	Δ mAP50–95	Δ mAP50	Δ Recall	× runtime
yolo26m	640	21.35	36.19	57.42	22.1	26.6	—	—	—	—
yolo26m	**1280**	**25.43**	**43.26**	**66.72**	72.2	75.6	**+4.08**	**+7.07**	**+9.30**	3.27
yolo26m	tiling 20 %	22.16	38.42	66.66	160.1	245.7	+0.81	+2.23	+9.24	7.24
yolo26x	640	23.21	39.01	60.41	47.9	51.6	—	—	—	—
yolo26x	**1280**	**27.02**	**45.52**	**70.24**	176.9	187.2	**+3.81**	**+6.51**	**+9.82**	3.70
yolo26x	tiling 20 %	24.09	41.26	70.02	388.1	617.8	+0.88	+2.25	+9.60	8.11



- **Recall** rises equally in both modes (+9–10 points): both find many more objects than 640.
- **mAP** rises mostly with 1280 (+4 points); with tiling by less than a point — the gain on small objects is eaten by the collapse on large ones (§5.2).
- **Runtime:** 1280 costs ×3.3–3.7, tiling ×7.2–8.1, and tiling's p95 is worse still (it depends on the number of tiles: 2 for a 960 px frame, 8 for 1920 px). On the plot, 1280 lies above and to the left of tiling for both models: **in accuracy per runtime, 1280 dominates tiling.**

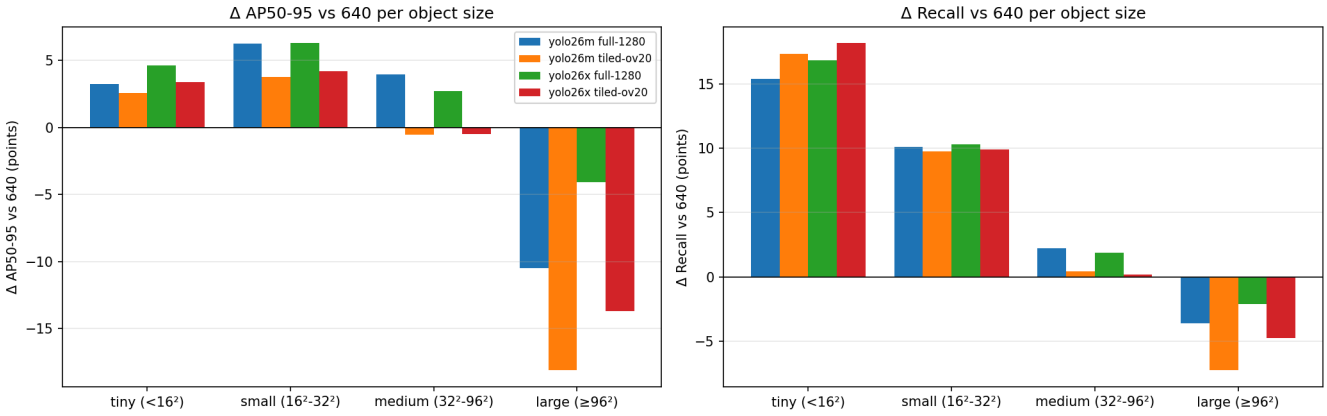
5.2. Per object size (Table 2)

yolo26m:

Size	Objects	Mode	AP50–95	AP50	Recall	Δ AP50–95	Δ AP50	Δ Recall
tiny (<16²)	27,345	640	6.21	15.31	25.29	—	—	—
		1280	**9.43**	**23.45**	40.68	**+3.22**	**+8.14**	+15.39
		tiling	8.79	21.75	**42.59**	+2.57	+6.43	**+17.30**
small (16²–32²)	23,491	640	15.64	32.22	65.48	—	—	—
		1280	**21.87**	**41.96**	**75.60**	**+6.23**	**+9.74**	**+10.12**
		tiling	19.39	37.96	75.24	+3.75	+5.74	+9.76
medium (32²–96²)	21,856	640	31.31	51.17	85.23	—	—	—
		1280	**35.25**	**55.54**	**87.46**	**+3.95**	**+4.37**	**+2.24**
		tiling	30.75	49.24	85.66	−0.56	−1.93	+0.43
large (≥96²)	2,409	640	**47.30**	**64.26**	**91.28**	—	—	—
		1280	36.82	52.52	87.67	−10.48	−11.74	−3.61
		tiling	29.19	42.55	84.02	**−18.11**	**−21.71**	**−7.26**

yolo26x:

Size	Objects	Mode	AP50–95	AP50	Recall	Δ AP50–95	Δ AP50	Δ Recall
tiny (<16²)	27,345	640	6.47	16.23	28.69	—	—	—
		1280	**11.08**	**26.34**	45.50	**+4.61**	**+10.11**	+16.82
		tiling	9.82	23.50	**46.83**	+3.36	+7.27	**+18.14**
small (16²–32²)	23,491	640	17.54	35.21	68.93	—	—	—
		1280	**23.80**	**44.93**	**79.25**	**+6.27**	**+9.72**	**+10.32**
		tiling	21.73	41.10	78.84	+4.19	+5.89	+9.91
medium (32²–96²)	21,856	640	34.22	54.82	87.44	—	—	—
		1280	**36.94**	**57.40**	**89.30**	**+2.72**	**+2.58**	**+1.86**
		tiling	33.74	53.36	87.61	−0.48	−1.46	+0.17
large (≥96²)	2,409	640	**44.98**	**60.45**	**92.36**	—	—	—
		1280	40.90	55.86	90.24	−4.08	−4.58	−2.12
		tiling	31.26	45.44	87.59	**−13.72**	**−15.00**	**−4.77**



Observations:

1. **Small objects (tiny + small = 68 % of all objects) are where the gain is.** Recall on tiny: 25 → 41–43 % (m) and 29 → 46–47 % (x). AP50 on tiny grows 1.4–1.6× — objects are actually *found*, not just outlined slightly better.
2. **AP50–95 on tiny objects stays low** (≤11) even in the best mode: found tiny objects are localised imprecisely (AP50 is 2.4–2.5× AP50–95). For a 10 px object, a 1–2 px error already means IoU < 0.75 — a property of the metric as much as of the detector.
3. **On tiny and small objects, 1280 beats tiling in AP** (by +0.6…+2.5 points), while tiling is only slightly ahead in tiny recall (+1.3…+1.9). The reason is §1: for 85 % of the frames, 1280 ≈ native resolution, so tiling has barely more pixels but cuts the image context.
4. **Medium:** 1280 helps (+2.7…+4.0), tiling does not (≈0).
5. **Large: both modes hurt**, tiling about twice as much (see §6).

5.3. Resolution vs model capacity (Table 3)

yolo26x − yolo26m within each mode, points:

	640	1280	tiling 20 %
Δ mAP50–95	+1.86	+1.59	+1.93
Δ mAP50	+2.82	+2.26	+2.83
Δ Recall	+2.99	+3.51	+3.35
Δ AP50–95 tiny	+0.26	+1.65	+1.04
Δ AP50–95 small	+1.90	+1.94	+2.35
Δ AP50–95 medium	+2.92	+1.69	+2.99
Δ AP50–95 large	−2.32	+4.08	+2.08
Δ Recall tiny	+3.40	+4.82	+4.24
× runtime (x / m)	2.16	2.45	2.42

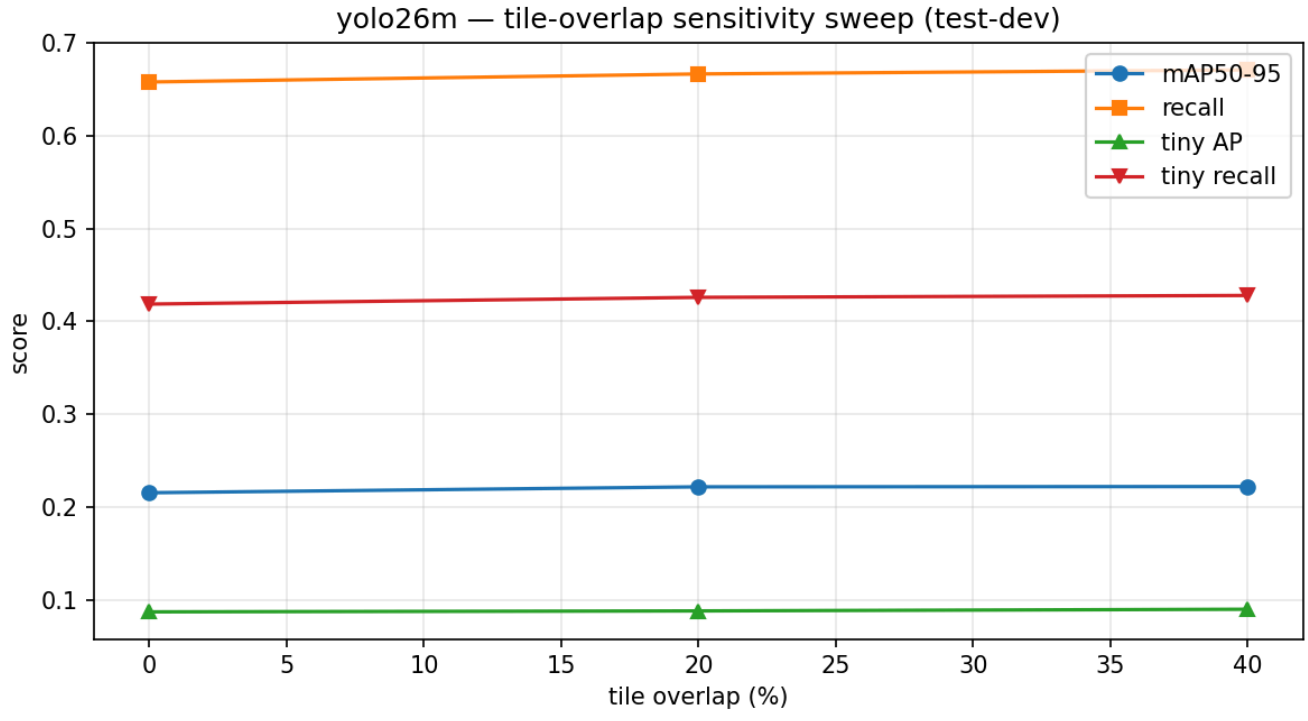
- The larger model gives **+1.6…1.9 mAP for ×2.2–2.5 runtime**. Raising the input to 1280 gives **+3.8…4.1 mAP for ×3.3–3.7**.
- **On tiny objects, capacity at 640 is almost useless (+0.26 AP):** once an object is squeezed to 3–5 px, the bigger model cannot see it either. At 1280 the same gap is +1.65 — capacity helps once the object is given pixels.
- At comparable cost: **yolo26m@1280 = 25.4 mAP in 72 ms** vs **yolo26x@640 = 23.2 mAP in 48 ms**. The best configuration overall is yolo26x@1280 (27.0 mAP, 177 ms).

Answer to the project's secondary question: for small objects, input resolution matters more than model capacity.

5.4. Tile overlap — sensitivity analysis (Table 4)

yolo26m only, everything else fixed. This is **not** hyperparameter tuning: the main value of 20 % was fixed in the proposal in advance; the sweep only shows how much the result depends on it (picking the best value on test-dev would be fitting to the test set).

Overlap	mAP50–95	Recall	AP tiny	AP small	AP medium	AP large	Recall tiny	ms (mean)	ms (p95)
0 %	21.53	65.80	8.69	19.06	29.83	28.11	41.86	149.7	184.8
20 %	22.16	66.66	8.79	19.39	30.75	29.19	42.59	160.1	245.7
40 %	22.20	67.10	8.97	19.50	30.64	29.43	42.79	200.6	469.7



- 0 → 20 %: +0.64 mAP, +0.86 recall, mostly on medium/large (objects on tile seams stop being lost), for +7 % mean runtime.
- 20 → 40 %: +0.04 mAP, +0.44 recall, while p95 runtime nearly doubles (245 → 470 ms): for wide frames the number of tiles grows from 8 to 15.
- Even 40 % overlap does not rescue large objects (AP large 29.4 vs 47.3 at 640): the problem is not the seams but the scale and the fact that the object is wider than a tile (§6).

5.5. Appendix: per native image size

The same metrics, separately for each group of frames (`results/metrics_by_native_size.csv`). The groups come from different cameras, altitudes and scenes, so compare **modes within a group**, not groups with each other. Δ is vs 640 within the same group; runtime is over the first 60 images of each group.

yolo26m:

Frame	Images	Scale 640 / 1280	Tiles	Δ mAP 1280 / tiling	Δ AP tiny 1280 / tiling	Δ Recall tiny 1280 / tiling	ms 640 / 1280 / tiling
960×540	49	×0.67 / ×1.33	2	+2.4 / +1.9	+3.8 / +4.0	+10.6 / +8.9	20 / 65 / 55
1360×765	167	×0.47 / ×0.94	6	+3.0 / −1.0	+3.0 / +2.6	+23.3 / +26.7	20 / 64 / 169
1400×788	933	×0.46 / ×0.91	6	+4.6 / −0.0	**+7.9 / +3.5**	+14.4 / +16.2	22 / 69 / 177
1400×1050	274	×0.46 / ×0.91	6	+4.7 / +1.5	+5.4 / +5.1	+17.0 / +19.1	27 / 88 / 173
1916×1078	151	×0.33 / ×0.67	8	+7.2 / +3.9	**+6.6 / +8.8**	+16.1 / +22.6	26 / 71 / 235
1920×1080	36	×0.33 / ×0.67	8	+4.0 / −0.3	**+3.3 / +8.9**	+27.4 / +34.7	25 / 70 / 241

yolo26x:

Frame	Images	Scale 640 / 1280	Tiles	Δ mAP 1280 / tiling	Δ AP tiny 1280 / tiling	Δ Recall tiny 1280 / tiling	ms 640 / 1280 / tiling
960×540	49	×0.67 / ×1.33	2	+1.9 / +0.2	+4.4 / +0.0	+14.3 / +10.9	45 / 164 / 122
1360×765	167	×0.47 / ×0.94	6	+3.7 / +0.8	+4.6 / +3.9	+24.0 / +27.5	46 / 162 / 442
1400×788	933	×0.46 / ×0.91	6	+5.3 / +1.2	+7.4 / +6.8	+14.7 / +16.4	48 / 165 / 472
1400×1050	274	×0.46 / ×0.91	6	+4.5 / +1.2	+6.1 / +5.1	+21.0 / +21.3	57 / 239 / 444
1916×1078	151	×0.33 / ×0.67	8	+7.3 / +4.9	**+6.4 / +7.5**	+19.3 / +22.6	51 / 170 / 599
1920×1080	36	×0.33 / ×0.67	8	+6.0 / +1.4	**+3.5 / +4.7**	+37.0 / +44.7	50 / 172 / 624

- **Where tiling really adds pixels (~1920 px frames: ×1.0 vs ×0.67 for 1280), it beats 1280 in AP on tiny objects** for both models (m: +8.8…+8.9 vs +3.3…+6.6). Where it does not (~1400 px, 85 % of frames: ×1.0 vs ×0.91), 1280 wins. So the hypothesis "tiling is best for the smallest objects" **holds where tiling has a resolution advantage** — only 12 % of test-dev frames.
- **1280 helps even on 960×540**, where the frame is only upscaled (+1.9…+2.4 mAP): the gain comes not only from new information but also from objects becoming larger in the detector's "units" (closer to the sizes it handles well).
- Tiling runtime grows with the number of tiles (2 → 8); 1280 runtime is nearly constant.
- The 960×540 (49 images, 1 large object) and 1920×1080 (36 images) groups are small, so their numbers are noisy; Δ on large for 960×540 is meaningless.

---

6. Why: scale mismatch and tiles

Train/test scale mismatch. The model was trained on frames downscaled to 640 (objects shrunk 1.5–3×) and learned the "usual" object sizes on its three output scales (stride 8/16/32). At 1280 inference, objects are ~2× larger than usual; with tiling, 1.5–3× larger:

- for **tiny/small** objects this helps: a 4 px object becomes an 8–12 px object the model has seen often;

- for **large** objects it hurts: a bus becomes larger than anything the model saw in training.

Evidence from the data: large objects also regress **at 1280, where there are no tiles at all** (−10.5 AP for m, −4.1 for x). So part of the large-object drop is the scale itself, not cutting. Tiling adds another −7.6 (m) and −9.6 (x) on top: an object wider than 640 px cannot fit into one tile, gets cut into pieces, and NMS merges the pieces poorly. Scale augmentation (±50 %) during training softens the mismatch but does not remove it. yolo26x's large-object drop at 1280 is half that of yolo26m — the bigger model is more robust to scale.

Why 1280 beats tiling on average: for 85 % of test-dev frames, 1280 ≈ native resolution, so tiling barely has more pixels, while 1280: (1) sees the whole frame — the model has context; (2) cuts nothing at seams; (3) has no duplicates on overlaps to merge; (4) needs one forward pass instead of 6–8.

---

7. Hypotheses check

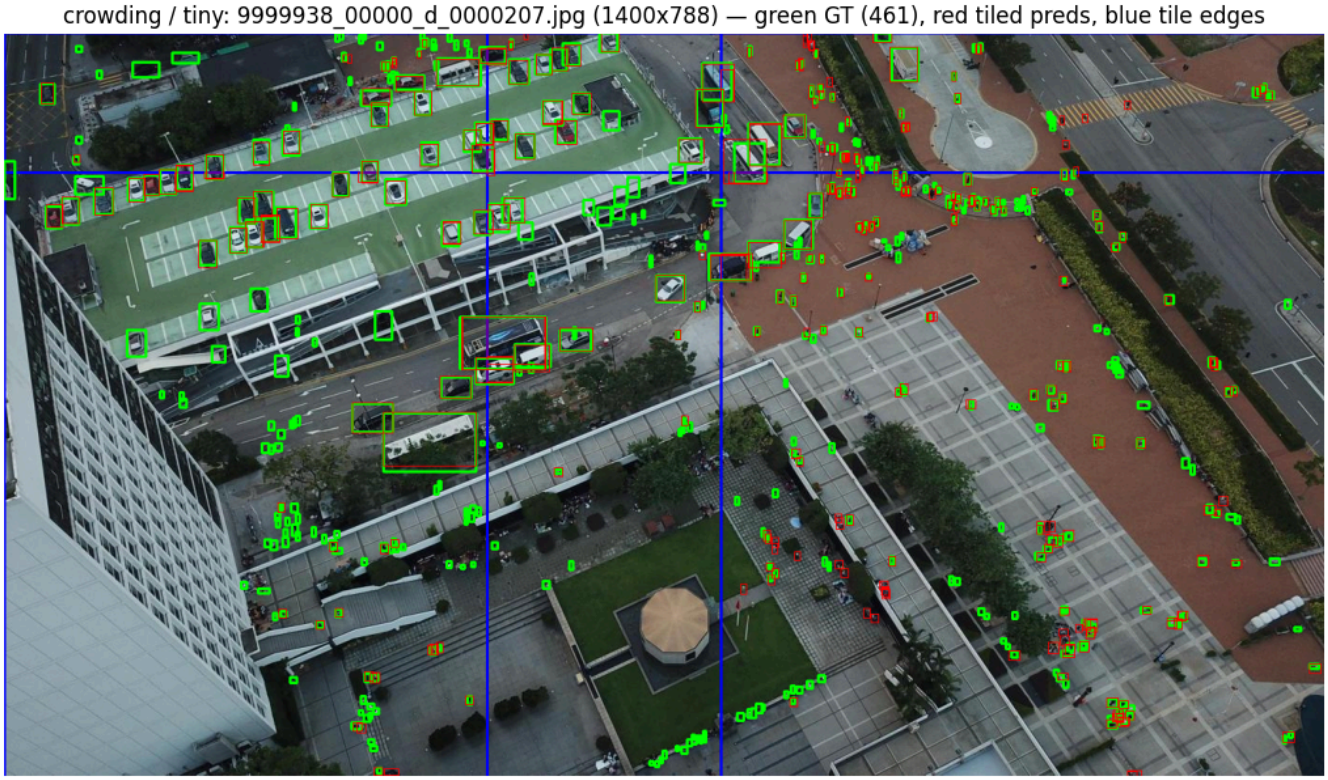
Hypothesis (proposal)	Verdict
Both tiling and 1280 raise small-object recall over 640	✅ **Confirmed.** Tiny recall +15…+18 points, small +10, for both modes and both models.
Tiling gains the most on the smallest objects	⚠️ **Partially.** In tiny recall tiling leads (+1.3…+1.9 over 1280); in tiny AP it trails (−0.6…−1.3). Tiling beats 1280 in tiny AP only on ~1920 px frames (§5.5), where it actually adds pixels.
Each gain comes at a measurable runtime cost	✅ **Confirmed.** 1280: ×3.3–3.7, tiling: ×7.2–8.1 (p95 up to ×12).
(secondary) Model capacity helps small objects as much as resolution	❌ **No.** x instead of m: +1.6…1.9 mAP; 1280 instead of 640: +3.8…4.1. On tiny objects at 640, capacity gives almost nothing (+0.26 AP).

---

8. Failure analysis

Three frames are selected deterministically, without cherry-picking (`tiling.ipynb`, section 5): the most crowded frame, the frame with the largest object, and the first frame where a large object crosses a tile boundary. Green = ground truth, red = tiling predictions with confidence ≥ 0.25, blue = tile boundaries; model yolo26m.

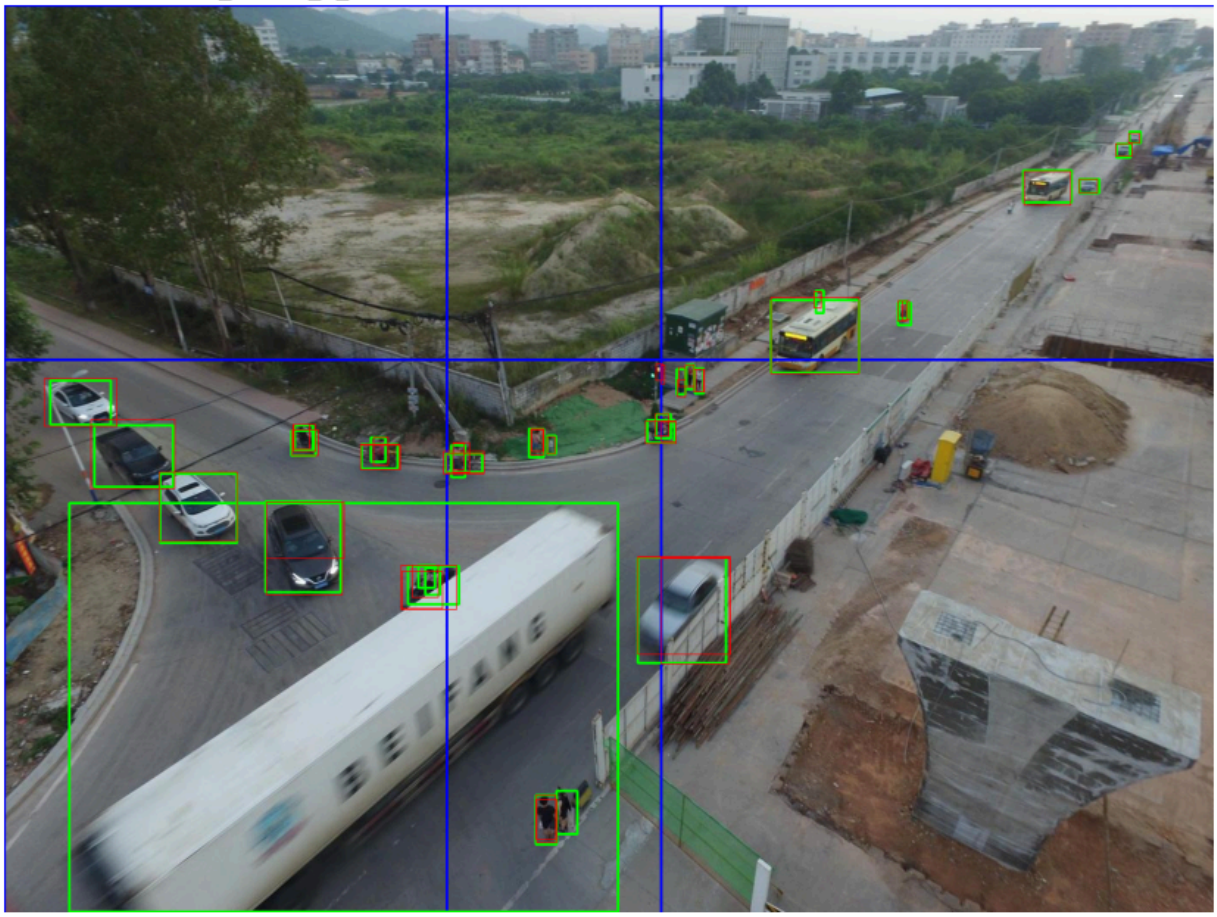
8.1. Crowds of tiny objects



`9999938_00000_d_0000207.jpg` (1400×788, **461 objects**). Cars in the parking lot and on the road are almost all found. The main misses are pedestrians on the square and on the sidewalk on the right — figures a few pixels high, often standing in dense groups — many green boxes without a red match. This is the dominant error type in every mode: even the best mode finds only ~41–47 % of tiny objects. This frame also shows why the shared detection cap was needed: on such frames the old cap of 300 boxes per image (but not per tile) gave tiling an unfair advantage.

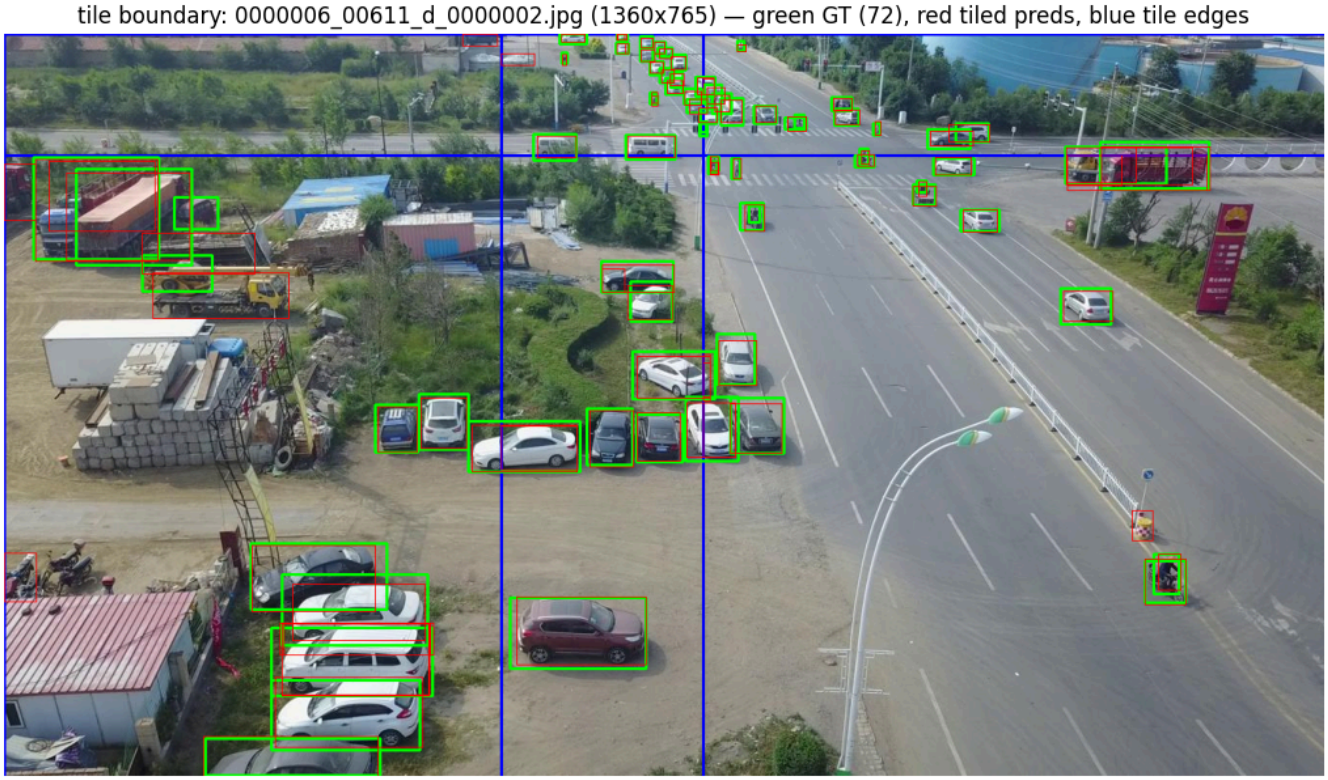
8.2. Large object cut by tiles (tiling regression)

large object: 9999986_00000_d_0000043.jpg (1400x1050) — green GT (32), red tiled preds, blue tile edges



`9999986_00000_d_0000043.jpg` (1400×1050). The truck in the lower left is wider than a tile and crosses two tile boundaries — **tiling does not detect it at all** (no red box), although at 640 such an object is an easy case. A car next to a tile edge gets a shifted box. This is the mechanism behind the large-object AP collapse (−18 for m): the object does not fit into one tile, and pieces from different tiles do not combine into one box.

8.3. Objects on tile boundaries



`0000006_00611_d_0000002.jpg` (1360×765). Trucks at the upper-left edge and in the upper right, crossing tile boundaries, get **several fragment boxes** instead of one; NMS does not merge them when the fragments overlap little. This produces both false positives (a fragment is not an object) and poor localisation. Cars fully inside a tile (frame centre) are outlined correctly. 20–40 % overlap mitigates this (§5.4) but does not eliminate it.

---

9. Limitations

1. **Training at 640 only** (by protocol) — hence the scale mismatch (§6). The comparison is fair for "what does changing the input do with the same weights", but does not answer "how to train best for each mode".
2. **Single run**, no confidence intervals or multiple seeds; differences < ~0.5 points are noise.
3. **Full train split instead of a "manageable subset"** from the task wording (the proposal justifies this: a subset would discard half the data; on Kaggle full training fit into ~8 h).
4. **Runtime measured on a Tesla T4**, batch 1, first 200 images; absolute numbers depend on the GPU, ratios between modes are more stable.
5. **val metrics in the training logs** come from the Ultralytics validator, not our evaluator.
6. **The 5-epoch stage-1 numbers** in `docs/baseline_results.md` / `docs/tiling_results.md` were computed before the detection cap was equalised — compare them with this report only approximately.
7. The 960×540 and 1920×1080 groups in the appendix are small (49 and 36 images).

---

10. Takeaways and next steps

- **Practical conclusion:** for small objects on VisDrone, the best trade-off is **full-image inference at 1280**; tiling pays off only for frames much larger than 1280 px and when recall on the very smallest objects is what matters.
- **Training matched to the inference mode:** training on 640×640 tiles at native resolution or multi-scale/1280 training would remove the scale mismatch and likely the large-object collapse.

- **Hybrid:** tiling + a full-image pass with merging (as in SAHI) — large objects from the full frame, small ones from the tiles.
- **Adaptive mode:** 1280 for frames ≤ 1400 px, tiling for ~1920 px frames (§5.5).

---

11. Reproduction

Step	Command / file	Output
Data	`python scripts/prepare_data.py`	`data/`
Training (Kaggle, 2× T4)	`bash scripts/kaggle_finetune.sh`	`runs/finetune/<model>_s640/weights/best.pt`
Evaluation (Kaggle, T4)	`bash scripts/kaggle_eval.sh`	`runs/{highres,tiling,results}/finetune/`
Result tables	`notebooks/results.ipynb`	`runs/results/finetune/results_tables.md`
Raw numbers	`results/metrics_all.csv`	every metric, 0–1
Archive	`python scripts/pack_results.py`	report + all computed results (+ checkpoints with `--with-weights`)

Training took 3.2 h (m) and 7.9 h (x) in parallel on two T4s; the full evaluation took 2.7 h on one T4 (~1.6 h for 640/1280 + appendix, ~1 h for tiling, the overlap sweep and failure cases).